

CS061 - Lab 1

Welcome to the LC-3!

1 High Level Description

Today's lab will get you set up, and cover the basics of how to use the LC3 Assembler/emulator, and have you write a tiny program in LC-3 Assembly language.

The specs for today's lab are rather lengthy, as there are a lot of new ideas & activities to cover. But don't panic - each step is pretty simple, and most future lab specs will be shorter (but harder!)

2 Objectives for This Week

1. Set up remote access to your CS account on the server sledge.cs.ucr.edu
2. Introduction to the basics of LC3 programming
3. Exercise 0: Hello World!
4. Exercise 1: Implement, run, and inspect a simple program yourself
5. So am I an assembly language programmer yet??

NOTICES:

1. ***We do encourage you to help one another and actively collaborate on your lab exercises (but NOT on your programming assignments, not EVER!)***

If you have taken this class before, you are required to delete all lab exercises and programming assignment code from your previous attempt.

See also the [class collaboration policy](#).

2. ***You will need a full-powered laptop or desktop computer to work on***

See the [BCOE laptop policy](#) for minimum laptop requirements.

3. ***Read "[Intro to Assembly Language](#)" (review of material covered in your first lecture).***

You MUST read this BEFORE lab.

In general, for all future labs, you must read the lab specs before attending lab, and make a start coding at least the first exercise before lab: this is even more important now that labs are virtual!

3 Lab 1

3.0 Getting set up

First, you need to set up remote access to your cs account from your local machine.

If you do not have a cs account yet, notify the TA immediately!

You will use **X2Go**, a package which provides remote connection to a server.

Download the appropriate X2Go client for your OS, following the instructions [here](#) (you must be logged into [R'Mail](#) to access this site).

Open the X2Go app, select Session->New Session.

In the New Session window, set:

Session name: to whatever you want (e.g. "sledge", or your name, or whatever)

Host: **sledge.cs.ucr.edu**

Login: your **NetID** (i.e. your UCR username, the name part of your @ucr.edu email address)

Session type: select **XFCE** from the drop-down list

Leave default values for all other settings.

When you click OK, this will create a session icon.

You will only need to do this part once.

Click the session icon to launch the connection: you will be asked to provide your cs account password.

Now [set up your cs061 directories](#) & get them ready for your labs & assignment.

When you are done with a remote session, do **NOT** end it by closing the X2Go window - your session will remain open & active on the server, chewing up resources and making the sysAdmin cry :(

ALWAYS end your session by logging out - click the little green running man icon at the bottom right of the screen.

3.1 Basic skeleton for LC-3 programs

Similar to C++, every LC-3 program has a certain amount of overhead necessary for making everything work properly.

In **C++**, you might have something like the following:

```
//  
// File: main.cpp  
// Description: A simple hello world-esque program  
// Author: Noobie Q. Programmer  
// Date created: 1/1/2019  
// Date last modified: 1/1/2019  
//  
#include <iostream>  
using namespace std;  
  
int main()  
{  
    cout << "Wow, this is the most fun I've had in years!" << endl;  
}
```

This source code is passed to a compiler, which ultimately produces a machine language executable.

In LC3, a basic program will have the following look:

```
;
; Author: Noobie Q. Programmer
; UCR NetID/CS username: nprog001
; Lab Section: xxx
; TA: Sean Foley
; Lab 1
; Date created: 1/1/2020
;

.orig x3000                ; **ALL** your programs will start at x3000
; Instructions (i.e. LC-3 code)
HALT                      ; end of program code
; Local Data                ; pseudo-ops for hard-coding data go here

.end                      ; .end is like the "}" after main() in C++.
                        ; It means "no more code or data"
```

If you haven't figured it out by now, LC-3 comments are denoted with semi-colons. They can be on a line by themselves or on the same line as actual code (at the end of the line).

This source code is passed to an assembler - a much simpler beast than a compiler, but like a compiler it produces a machine language executable.

Pseudo-ops (*the keywords that start with a '.' - see next section for more details*) are like compiler directives, in that they tell the assembler how to set things up.

3.2 Basics of LC-3 programming

An LC-3 program consists of three basic elements: **instructions**, **pseudo-ops**, and **labels**.

Labels are simply "aliases" for memory addresses. These are really, really useful because they relieve the programmer of the responsibility of figuring out which line of code is at which memory address, allowing us to use symbolic names to point to data and instructions.

Pseudo-ops tell the assembler how to set things up before starting to translate the source code into machine language. They are a bit like compiler directives in C++, such as `#include`.

For example, the pseudo-op **.FILL** tells the assembler to write a hard-coded value into memory (i.e. RAM, or system memory). For example, the two **.FILL** pseudo-ops in Table 1 would write the decimal value 6 into memory at x3005 (*aka DEC6*), and the decimal value 12 into memory at x3006 (*aka DEC12*) before the program begins execution.

Note that you just need to label the lines - you don't need to know the actual memory addresses corresponding to those labels (that's the assembler's job!).

Table 1: Labels and Pseudo-ops

<u>Memory address</u>	<u>Label</u>	<u>Pseudo-Op</u>	<u>Hard-coded value</u>
x3005	DEC6	.FILL	#6
x3006	DEC12	.FILL	#12

(The '#' before the numbers means decimal, i.e base 10; 'x' means hexadecimal, i.e. base 16).

Now we can refer to these memory locations (and therefore access the data stored there) using the names DEC6 and DEC12.

We will learn about another pseudo-op **.STRINGZ** in a moment (in section 3.3).

Instructions: There are 15 LC-3 instructions, in three distinct categories (see [this tutorial](#) for more):

- Arithmetic / Logic operations (addition, bitwise AND, bitwise NOT)
- Data movement (moving data back & forth between memory and local registers)
- Program control (changing the order of execution of code - if statements, for loops, etc)

The following sections illustrate one or two instructions from each of these categories.

3.2.1 Arithmetic / Logic Operations

These instructions are used for manipulating values and doing calculations - Adding two numbers, bitwise ANDing two words together, or taking the bitwise logical NOT of a word.

The LC-3 has 8 general purpose, 16-bit **local registers** (analogous to variables in C++), named R0 through R7, that can be used as temporary storage for performing such operations.

The **ADD** instruction (2's complement integer data type)

This instruction adds two numbers, and writes the result to a register. It comes in two flavors:

One in which both operands come from local registers; the other in which the first operand comes from a register, while the second is actually embedded in the instruction itself ("immediate" mode).

In the immediate mode, the embedded value can only be in the range {-16, +15}.

See the [ADD tutorial](#) (also on the *Piazza Resources* page, "LC3 instruction set") for more details

3.2.2 Data movement

Data movement instructions are used when you want to copy a value from a memory location into a local register ("loading"), or copy a value from a register to a memory location ("storing").

You will be doing this a **lot**, since the LC3 can operate only on values in registers, not in memory.

The **LD** instruction ("Load Direct" - probably the most frequently used instruction!)

This instruction copies the contents of a specified memory location (identified by a label) into a register.

See the [LD tutorial](#)

The **LEA** instruction ("Load Effective Address")

This instruction translates a label – the name you give to a memory location – back into the 16-bit memory address it stands for, and stores that address into the specified register.

There is no actual data movement here – just the decoding of a label.

See the [LEA tutorial](#)

3.2.3 Program control

These instructions are used in control structures such as **BRANCHING**, **LOOPING**, and **SUBROUTINE CALLS** (a simpler version of function calls in higher-level languages).

By default, a program starts at the first instruction and executes one instruction after another in order until it reaches the end ("sequential execution").

Control instructions allow us to alter the default sequencing in several ways.

The **BR** instruction ("Conditional Branch") can be used to transfer control to the instruction at a specified label based on a *condition*.

All the control structures you are familiar with from C++ (if; if-else; if-else if-else; for; while; do-while; switch) can be built using this instruction, sometimes in combination with the **JMP** instruction.

But before we can learn how it works, we have to know about the [Condition Codes](#):

Three "flags" (single-bit registers) are set whenever a value is written into any one of the local registers, indicating whether the value being written is Negative, Zero, or Positive. These flags are called the **N Z P Condition Codes**, and allow us to make decisions based on the last value written to a register - referred to as the last modified register ("**LMR**").

The decisions are made by the conditional branch instruction BR, which has three possible modifiers: {**n, z, p**}.

As you can guess, the n modifier asks whether the N Condition Code is currently set or not (*was the value just written negative?*); and likewise for the z and p modifiers.

The BR instruction can be modified by any combination of n, z, p. It will cause a branch to the labelled instruction **IF AND ONLY IF** any one of the specified conditions is met.

See the [BR tutorial](#)

And now, we can finally get started on the actual lab exercises :)

3.3 Exercise 0: lab1_ex0 (aaargh!! not “hello world” again?!?!)

As always, the very first program we will write will simply output the message “Hello World!” Though simple, it will illustrate several of the points above, plus some things you will fully understand only later on (*exactly like when you first encountered cout in C++*):

- We will store the string as an array of ASCII characters in memory, starting at an address labelled “message”, using the pseudo-op **.STRINGZ**
This pseudo-op stores a string in memory, one character per memory address, with a zero (“the null character” or ‘\0’) after the last character.
You will recognize this as a “c-string”, i.e. a null-terminated character array. So the line:
 shrug .STRINGZ "whatever"
will write the characters 'w' 'h' 'a' 't' 'e' 'v' 'e' 'r' followed by x0000, to the nine memory locations starting with the labelled address shrug.
- We will use the LEA instruction to translate the label "shrug" into its actual address, which we will store in R0.
- We will invoke an i/o subroutine called PUTS (*can also be written as TRAP x22*), which uses R0 to locate the string (*it knows when to stop because the last value in the char array is '\0'*).

```
;
; Hello world example program
; Also illustrates how to use PUTS (aka: Trap x22)
;
; .ORIG x3000
; -----
; Instructions
; -----
;      LEA R0, MSG_TO_PRINT    ; R0 <-- the location of the label: MSG_TO_PRINT
;      PUTS                    ; Prints string defined at MSG_TO_PRINT
;
;      HALT                    ; terminate program
; -----
; Local data
; -----
;      MSG_TO_PRINT    .STRINGZ    "Hello world!!!\n"    ; store 'H' in an address labelled
;                                                           ; MSG_TO_PRINT and then each
;                                                           ; character ('e', 'l', 'l', 'o', ' ', 'w', ...) in
;                                                           ; it's own (consecutive) memory
;                                                           ; address, followed by #0 at the end
;                                                           ; of the string to mark the end of the
;                                                           ; string
; .END
```

Now write this program for yourself:

1. In your terminal, navigate to cs061/labs/lab1, and open your exercise 0 file in geany:
 geany lab1_ex0.asm &
2. Copy in the code from the above image & save the file (*you can leave the editor open if you wish*).
3. From your terminal, launch the LC-3 assembler/emulator:
 simpl lab1_ex0.asm &

This will open two windows:

- a. the **emulator** itself, showing the state of the processor, including the registers, the condition codes, the assembled code ("machine language" or "binary") in memory; and
 - b. the **console**, where all output from the program is displayed, and where you can type input to the program. *(You may have to move the emulator to find the console!)*
4. Now you will open a third, the **Text Window**. This can be used to enter command line input to the emulator, but we will use it mostly to watch for warnings and error messages.
Check the Text Window box at the bottom center of the emulator window.
5. Hit the **Run** button at the bottom of the emulator window (or enter the word "run" in the Text Window input box), and observe the words "Hello World!" magically appear in the console window!

3.4 Exercise 1: lab1_ex1 - A "real" program

Now Hello World is out of the way, we can write a program that actually does some processing: it will multiply a number, which we will store in advance in memory, by six *(How hard can that be, right??!?)*

We will also take this opportunity to introduce the style you **MUST** use for all your LC-3 programs.

See the image below:

First comes the header, which will always be included in your starter code: obviously, you have to always fill it out with your personal details.

Next, we have the line `.orig x3000` *(you may use upper or lower case for all LC-3 instructions, as you wish)*. This indicates where in memory the code will reside. The LC3 assembler requires all programs to be loaded at or above x3000: all your programs **must** load the "main" routine to x3000 *(just try something else and see what happens!! On second thoughts – don't!)*

The code itself multiplies 6 by 12.

Since the LC-3 has no multiply instruction *(really!!)*, the operation has to be carried out by adding 12 into an accumulator 6 times.

We could have added 6 into an accumulator 12 times: why would that not be a good idea?

Can you see how this program works?

The next three lines load hard-coded values from memory into registers R1, R2, and R3 respectively:

R1 is an "accumulator" - i.e. at the end, it will hold the result - so we have to initialize it to 0;

R2 holds the multiplicand (*the number to be multiplied*); and

R3 holds the multiplier - we will use it as a loop counter, to keep track of how many times the loop will execute (*i.e. we will create a "counter-controlled" loop*).

Make sure you understand how the BRp instruction controls the loop - this is the foundation of all Assembly Language control structures!!

TIP:

Actually, copying values from memory to register (and vice versa) is rather inefficient.

Mostly, we can't avoid it - our fixed data has to be stored in memory.

However, there is one special value for which we can do better - 0!

There is a much cleaner way to "zero out" a register than copying a hard coded 0 from memory to register:

AND R1, R1, x0 ; R1 <- (R1) AND x0000

This does what we call a "bitwise AND" of R1 with 0, i.e. it ANDs each of the 16 bits in R1 with the corresponding bit in x0000, and stores the result (x0000) back in R1 - a LOT quicker than reading from memory!

Note how the program data is stored in a block separated from the code, **following** the HALT instruction, which ends code execution - this way, the processor does not attempt to "execute" the data as though it were code!

And note how we use comments to clearly distinguish the two blocks.

Now that you have had a nice lengthy overview of some basic LC-3 instructions and have seen how to use the simpl assembler/emulator, it is time for you to implement the program yourself.

Open file **lab1_ex1.asm** for editing, type in the code from the image, save it, and open it in simpl. Remember to replace the first LD R1 instruction with the AND instruction, as per the tip above.

This time, instead of hitting Run, step through each line of code using the **Step** button in the emulator until you reach the HALT instruction. Notice how the value of R1 keeps going up by 12 until it reaches (guess what??) 72
(Now, how did it know to stop at just the right number??)

Finally, congratulate yourself on a job well done :)

3.5 Submission

Do not submit lab 1 - instead, just demo your work to your TA **before you leave lab**.

Your TA will maintain a Google doc where you can signal that you're ready to demo (and another to request assistance)

In later labs, you may occasionally get stuck and not be able to complete all exercises in the 3 hours: in that case, you must still show your TA how far you got, and request permission to complete it after lab.

Your TA will usually give you partial credit for what you have done (*so long as you have actually been present for the whole time*), and allow you to complete & demo the rest later for full credit, so long as you have worked at it seriously.

When you complete the exercises, demo it to either of the TAs in their office hours **before your next lab**. Office hours are posted on Piazza, under the "Staff" tab.

4 So what do I know now?

You have now mastered the following skills:

- How to write basic LC-3 assembly language code
- What labels are and how to use them
- How to use the .FILL and .STRINGZ pseudo-ops to store program data in memory
- The LD, LEA, BR, HALT (aka Trap x25) instructions
- The PUTS output routine (aka: Trap x22)
- How to create, save, run, and step through an LC-3 program
- And you now know that the LC-3 can't even multiply two integers by itself!

Not bad for the first day :)

